

CS291A: Scalable Internet Services

gTrack: Track Prices of Games on Steam

Nazmus Saquib Udit Paul Alex Ermakov

Graduate Students
Department of Computer Science
University of California Santa Barbara

December 8, 2018



Outline

- 1 Introduction
- 2 Data Model
- 3 Motivation
- 4 setup
- 5 Results
- 6 r^2
- 7 r^3
- 8 r^4
- 9 ajax
- 10 index
- 11 caching
- 12 Scaling
- 13 16
- 14 32
- 15 64
- 16 conclusions

Outline

- 1 Introduction
- 2 Data Model
- 3 Motivation
- 4 setup
- 5 Results
- 6 r_2
- 7 r_3
- 8 r_4
- 9 ajax
- 10 index
- 11 caching
- 12 Scaling
- 13 16
- 14 32
- 15 64
- 16 conclusions

Introduction

- 1 gTrack is a website built using Ruby on Rails.
- 2 gTrack is designed for users to get information related to the games available on Steam.
- 3 Logged in users can comment and express their like or dislike about any game.
- 4 gTrack users are presented with a highly specialized search feature.

Outline

- 1 Introduction
- 2 Data Model**
- 3 Motivation
- 4 setup
- 5 Results
- 6 r^2
- 7 r^3
- 8 r^4
- 9 ajax
- 10 index
- 11 caching
- 12 Scaling
- 13 16
- 14 32
- 15 64
- 16 conclusions

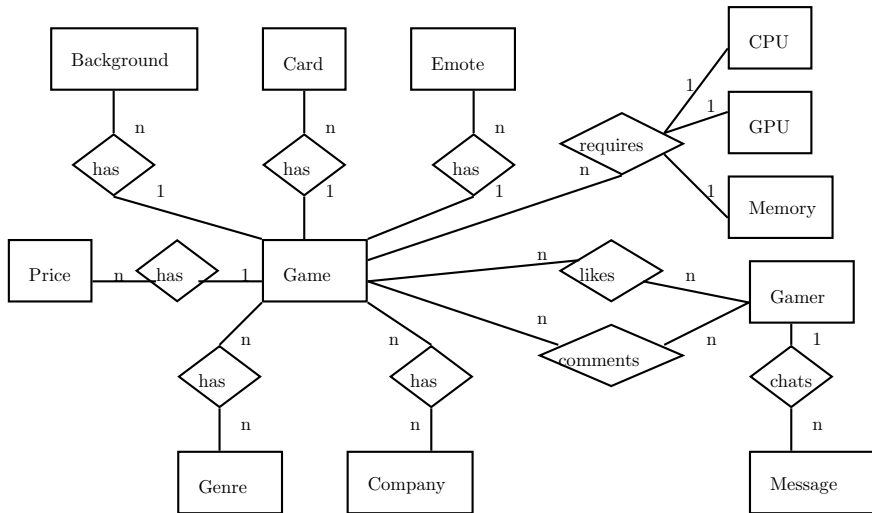
Motivation

- 1 Steam is the largest PC game distribution platform, yet its search functionalities are inadequate in meeting specialized queries.
- 2 Items such as emotes, cards and background to a game are not presented in an organised manner in Steam.
- 3 The games available on Steam do not have their price histories.

Outline

- 1 Introduction
- 2 Data Model
- 3 Motivation**
- 4 setup
- 5 Results
- 6 r^2
- 7 r^3
- 8 r^4
- 9 ajax
- 10 index
- 11 caching
- 12 Scaling
- 13 16
- 14 32
- 15 64
- 16 conclusions

Entity Relationship Diagram



Overview of Seed Data

- In total **349 MB** worth of data
- Major tables:
 - 15450 games
 - 775510 comments (50 comments/game on average)
 - 436322 price history (28 histories/game on average)
 - 3095 backgrounds, 9391 cards, 3662 emotes

Outline

- 1 Introduction
- 2 Data Model
- 3 Motivation
- 4 setup**
- 5 Results
- 6 r^2
- 7 r^3
- 8 r^4
- 9 ajax
- 10 index
- 11 caching
- 12 Scaling
- 13 16
- 14 32
- 15 64
- 16 conclusions

Test set-up

- ① User arrival rates were modelled in 8 phases.
- ② The work flow consisted of 4 distinct sessions with various probabilities.
- ③ Interspersed waiting within sessions.
- ④ Specialized tests were set up to test caching.

Outline

- 1 Introduction
- 2 Data Model
- 3 Motivation
- 4 setup
- 5 Results**
- 6 r^2
- 7 r^3
- 8 r^4
- 9 ajax
- 10 index
- 11 caching
- 12 Scaling
- 13 16
- 14 32
- 15 64
- 16 conclusions

Results

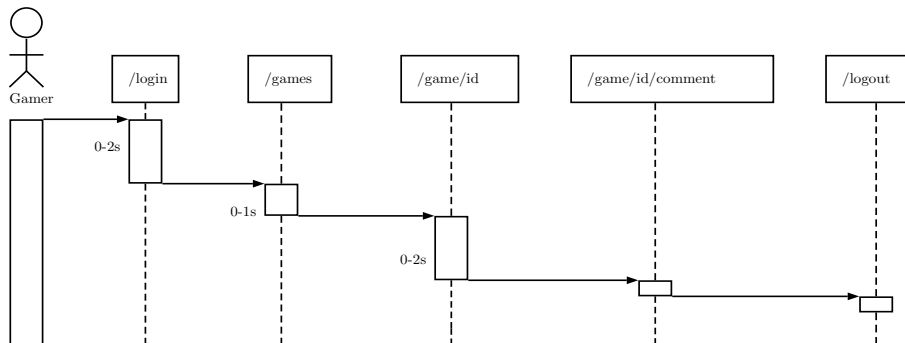


Figure: First Session

Outline

- 1 Introduction
- 2 Data Model
- 3 Motivation
- 4 setup
- 5 Results
- 6 r2**
- 7 r3
- 8 r4
- 9 ajax
- 10 index
- 11 caching
- 12 Scaling
- 13 16
- 14 32
- 15 64
- 16 conclusions

Session 2

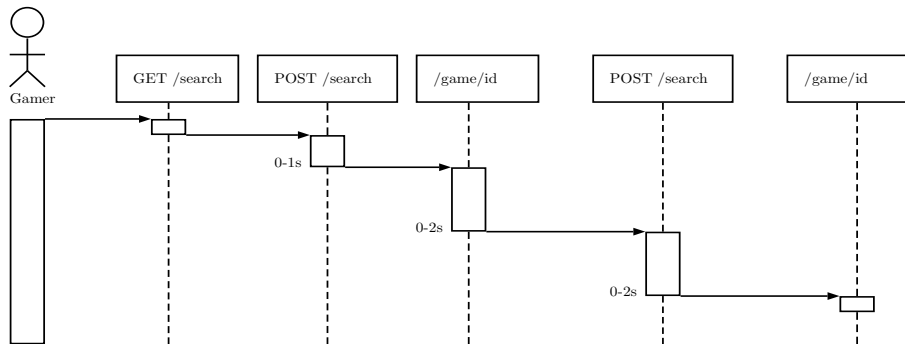


Figure: Second Session

Outline

- 1 Introduction
- 2 Data Model
- 3 Motivation
- 4 setup
- 5 Results
- 6 r2
- 7 r3**
- 8 r4
- 9 ajax
- 10 index
- 11 caching
- 12 Scaling
- 13 16
- 14 32
- 15 64
- 16 conclusions

Session 3




Figure: Third Session

Outline

- 1 Introduction
- 2 Data Model
- 3 Motivation
- 4 setup
- 5 Results
- 6 r^2
- 7 r^3
- 8 r^4**
- 9 ajax
- 10 index
- 11 caching
- 12 Scaling
- 13 16
- 14 32
- 15 64
- 16 conclusions

Session 4





Figure: Fourth Session

Outline

- 1 Introduction
- 2 Data Model
- 3 Motivation
- 4 setup
- 5 Results
- 6 r_2
- 7 r_3
- 8 r_4
- 9 ajax**
- 10 index
- 11 caching
- 12 Scaling
- 13 16
- 14 32
- 15 64
- 16 conclusions

Optimization 1: AJAX



Outline

- 1 Introduction
- 2 Data Model
- 3 Motivation
- 4 setup
- 5 Results
- 6 r^2
- 7 r^3
- 8 r^4
- 9 ajax
- 10 index**
- 11 caching
- 12 Scaling
- 13 16
- 14 32
- 15 64
- 16 conclusions

Optimization 2: Indexing




Figure: Mean duration for index page transaction without and with indexing.

Outline

- 1 Introduction
- 2 Data Model
- 3 Motivation
- 4 setup
- 5 Results
- 6 r_2
- 7 r_3
- 8 r_4
- 9 ajax
- 10 index
- 11 caching**
- 12 Scaling
- 13 16
- 14 32
- 15 64
- 16 conclusions

Optimization 3: Caching




Figure: Mean duration for system requirement search transaction with and without caching.

Outline

- 1 Introduction
- 2 Data Model
- 3 Motivation
- 4 setup
- 5 Results
- 6 r^2
- 7 r^3
- 8 r^4
- 9 ajax
- 10 index
- 11 caching
- 12 Scaling**
- 13 16
- 14 32
- 15 64
- 16 conclusions

Horizontal and Vertical Scaling

- 1 The website was load tested with various hardware configuration.
- 2 It was detected very early that the major bottleneck lay with the database.
- 3 The app server used was c5 with various database servers.

Outline

- 1 Introduction
- 2 Data Model
- 3 Motivation
- 4 setup
- 5 Results
- 6 r^2
- 7 r^3
- 8 r^4
- 9 ajax
- 10 index
- 11 caching
- 12 Scaling
- 13 16**
- 14 32
- 15 64
- 16 conclusions

16 users/second arrival rate




Figure: Mean response time while handling up to 16 users/second

Outline

- 1 Introduction
- 2 Data Model
- 3 Motivation
- 4 setup
- 5 Results
- 6 r^2
- 7 r^3
- 8 r^4
- 9 ajax
- 10 index
- 11 caching
- 12 Scaling
- 13 16
- 14 32**
- 15 64
- 16 conclusions

32 users/second arrival rate

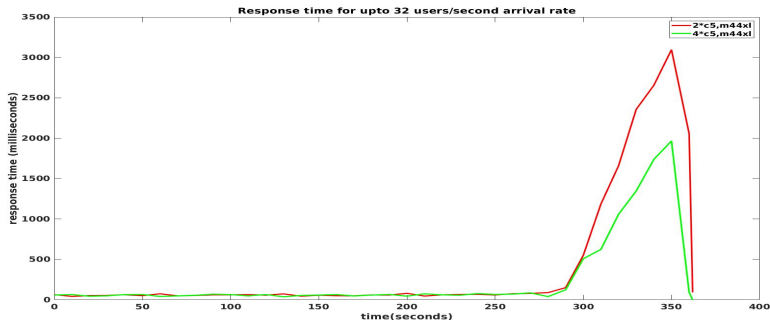


Figure: Mean response time while handling up to 32 users/second

Outline

- 1 Introduction
- 2 Data Model
- 3 Motivation
- 4 setup
- 5 Results
- 6 r_2
- 7 r_3
- 8 r_4
- 9 ajax
- 10 index
- 11 caching
- 12 Scaling
- 13 16
- 14 32
- 15 64**
- 16 conclusions

64 users/second arrival rate




Figure: Mean response time while handling up to 64 users/second

Outline

- 1 Introduction
- 2 Data Model
- 3 Motivation
- 4 setup
- 5 Results
- 6 r_2
- 7 r_3
- 8 r_4
- 9 ajax
- 10 index
- 11 caching
- 12 Scaling
- 13 16
- 14 32
- 15 64
- 16 conclusions**

Conclusion

1